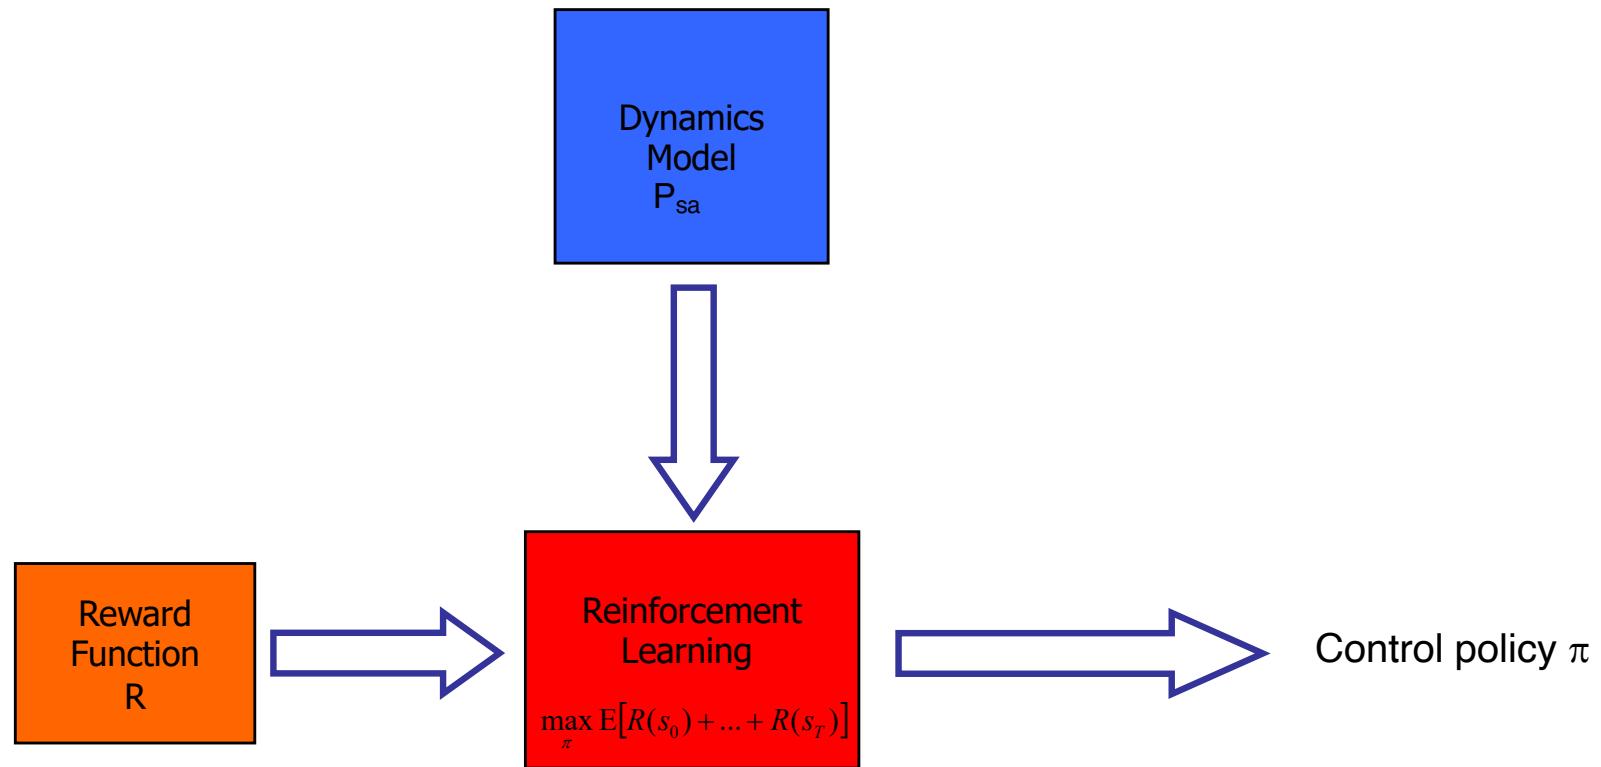


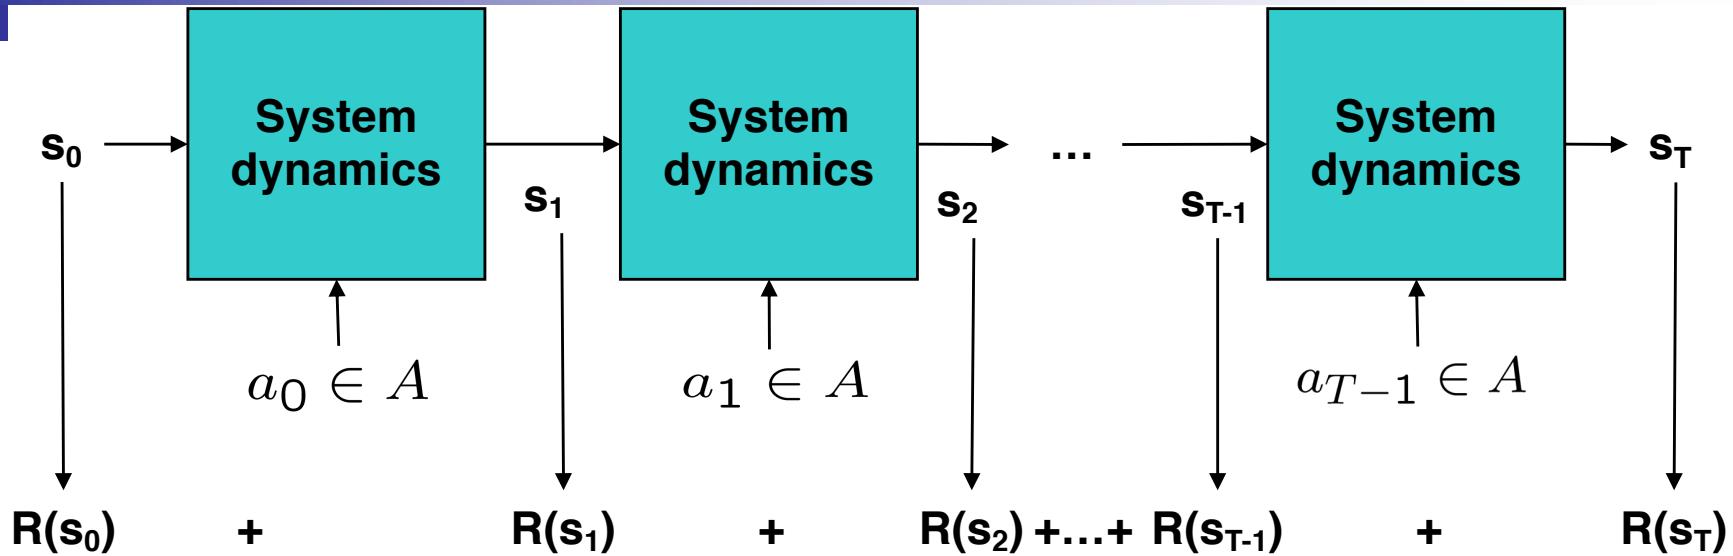
# **Learning by Demonstration: Imitation Learning, Inverse Reinforcement Learning, and Apprenticeship Learning**

Luís Macedo  
University of Coimbra

# RL formalism



# RL formalism

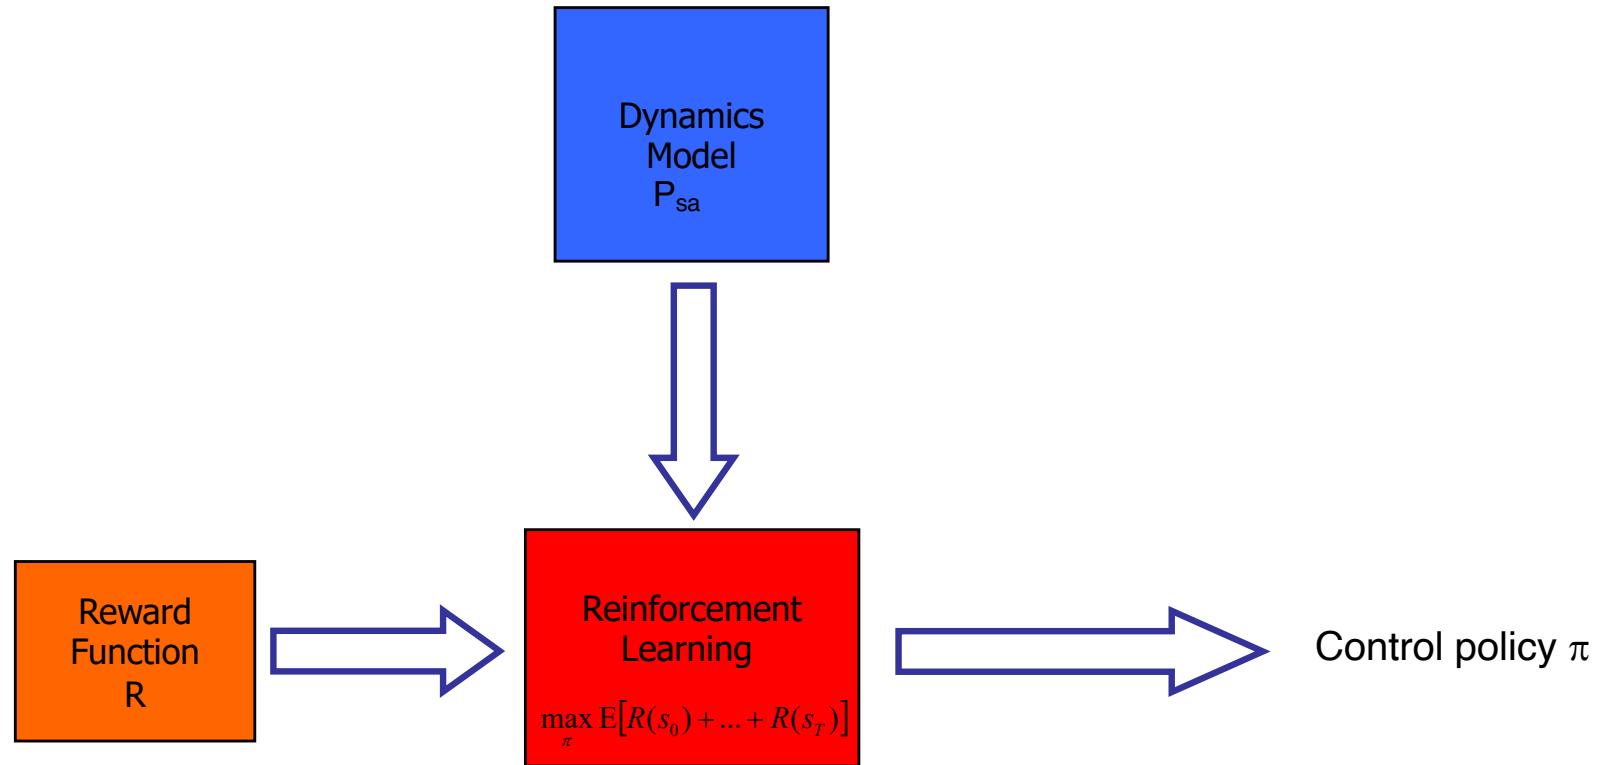


- Assume that at each time step, our system is in some state  $s_t$ .
- Upon taking an action  $a$ , our state randomly transitions to some new state  $s_{t+1}$ .
- We are also given a reward function  $R$ .
- The goal: Pick actions over time so as to maximize the expected score:  $E[R(s_0) + R(s_1) + \dots + R(s_T)]$ .

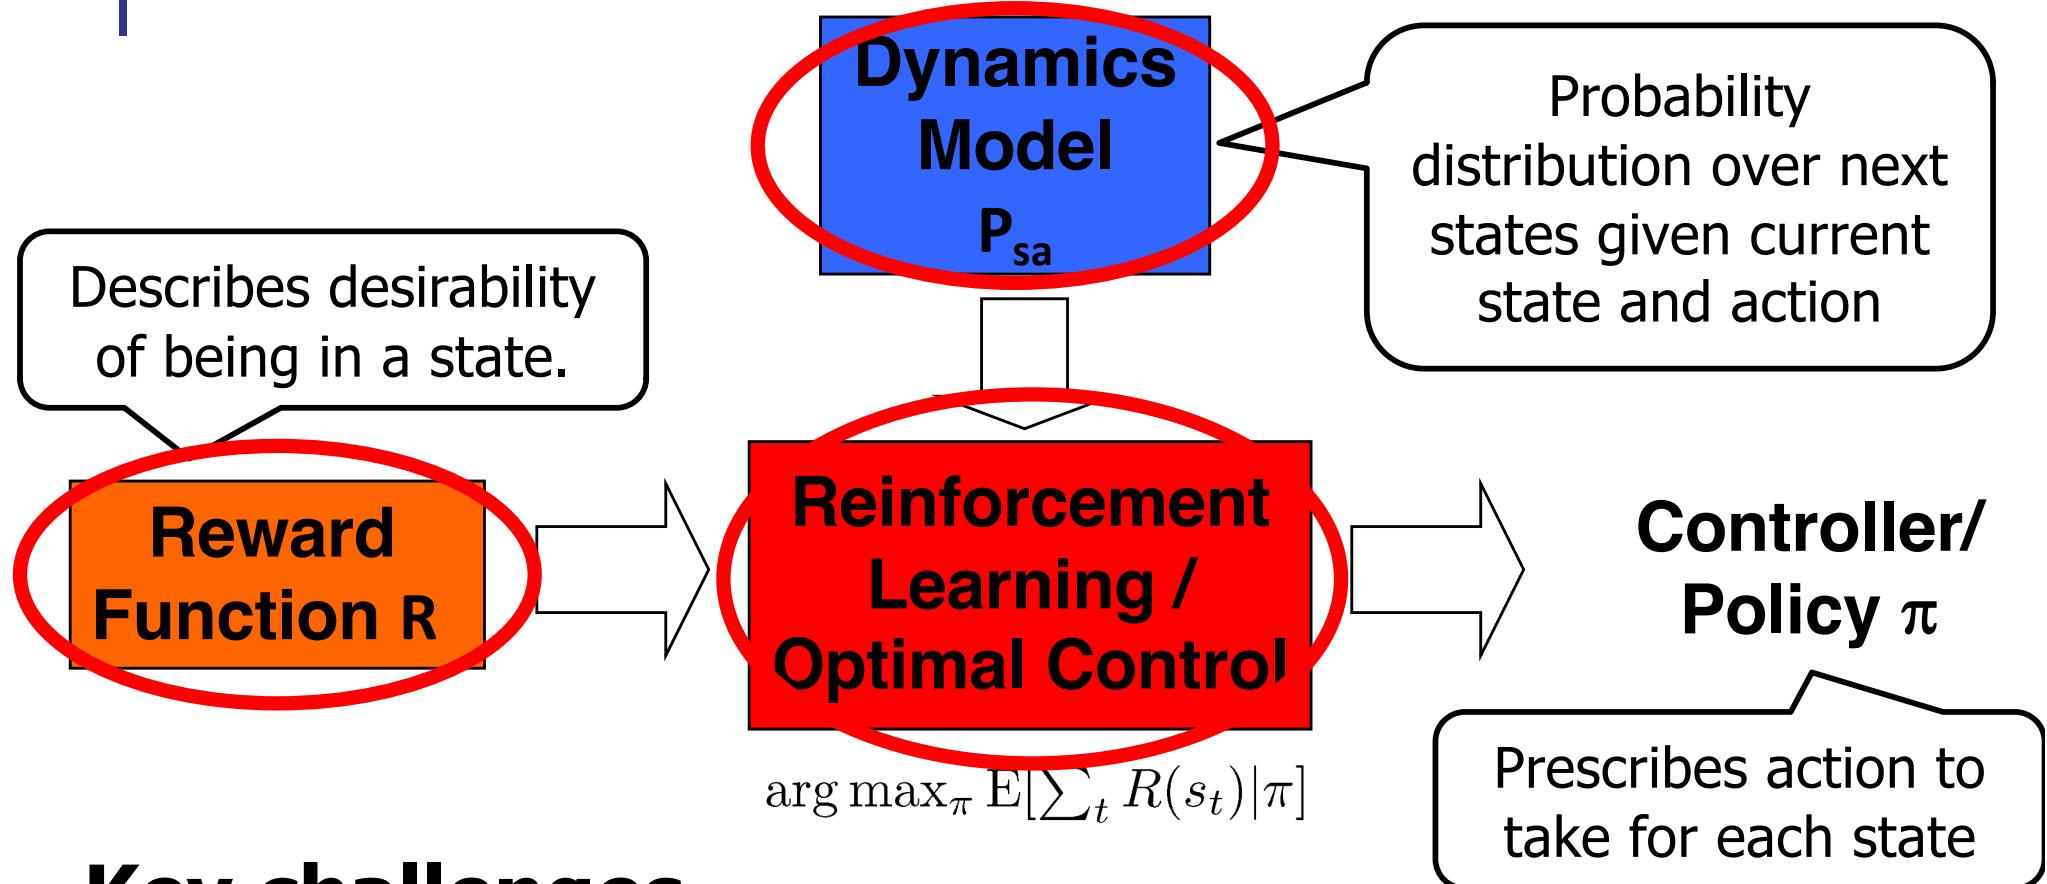
# RL formalism

- Markov Decision Process  $(S, A, P_{sa}, s_0, R)$   
$$R(s) = w^T \phi(s),$$
  
$$\phi : S \rightarrow [0, 1]^k$$
:  $k$ -dimensional feature vector.  
$$\|w\|_2 \leq 1.$$
- W.l.o.g. we assume
- Policy  $\pi : S \rightarrow A.$
- Utility of a policy  $\pi$  for reward  $R = w^T \phi$
- $$U_w(\pi) = E[\sum_{t=0}^T R(s_t) | \pi]$$

# RL formalism



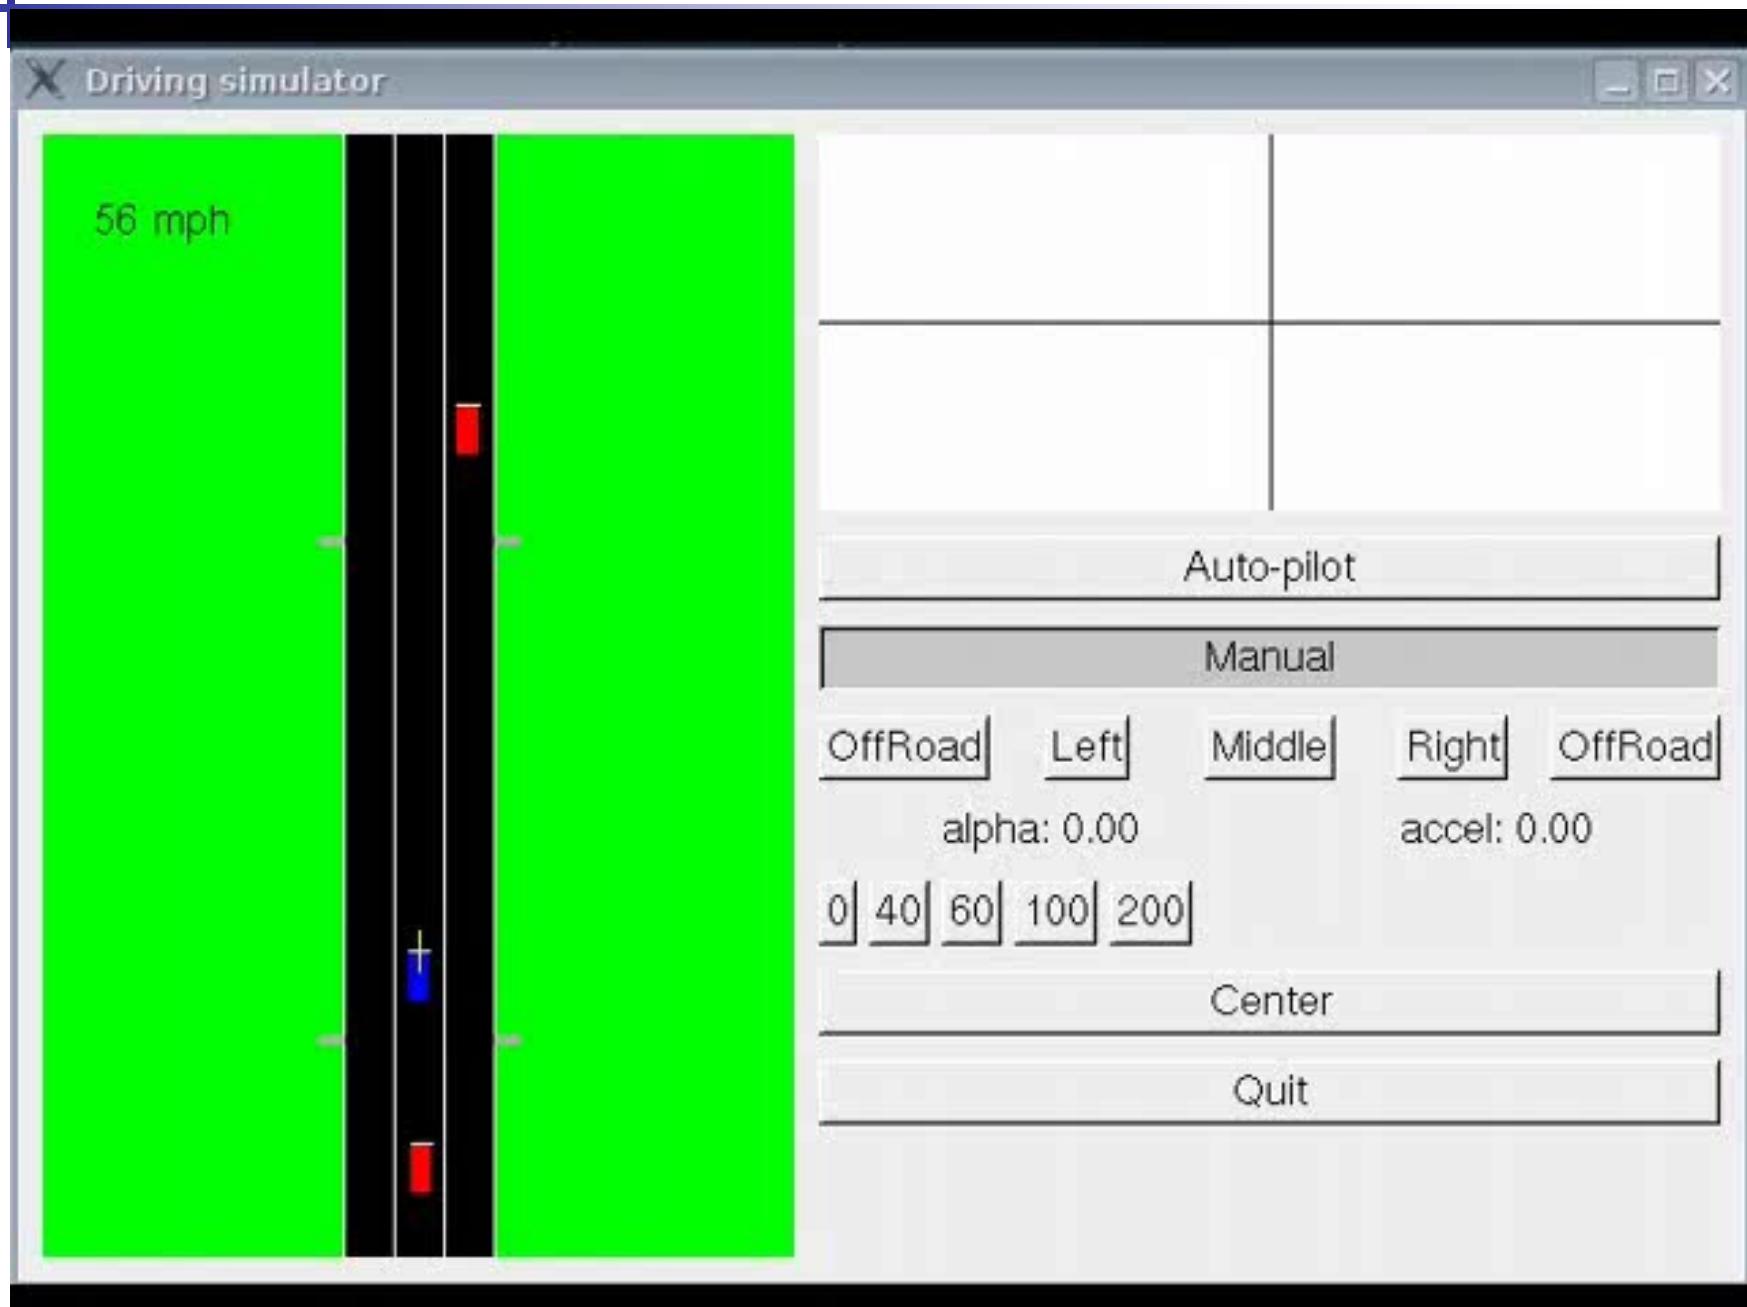
# Big picture and key challenges



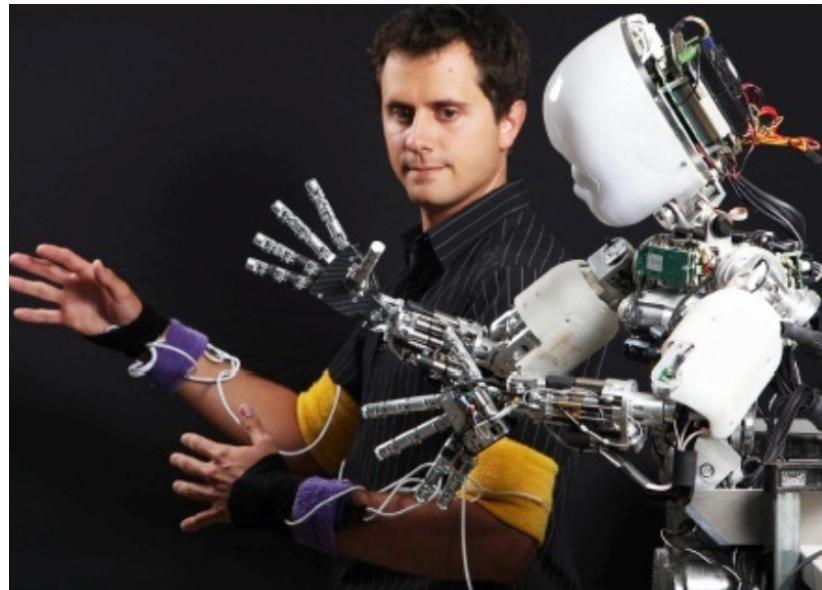
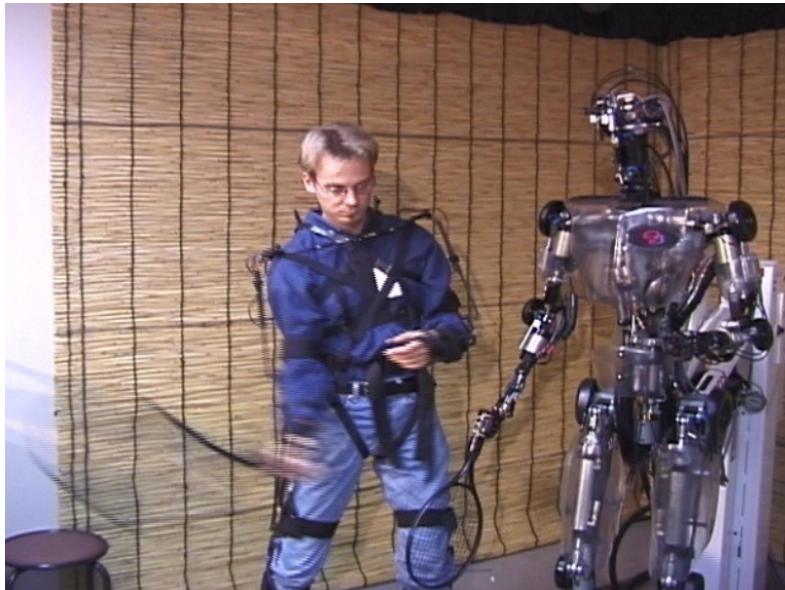
## ■ Key challenges

- Providing a formal specification of the control task.
- Building a good dynamics model.
- Finding closed-loop controllers.

# Example task: driving



# Imitation Learning



# Problem setup

- Input:
  - Dynamics model / Simulator  $P_{sa}(s_{t+1} | s_t, a_t)$
  - *No* reward function
  - Teacher's demonstration:  $s_0, a_0, s_1, a_1, s_2, a_2, \dots$   
(= trace of the teacher's policy  $\pi^*$ )
- Desired output:
  - Policy  $\pi : S \rightarrow A$ , which (ideally) has performance guarantees, i.e.,
$$\mathbb{E}\left[\frac{1}{T} \sum_t R^*(s_t) | \pi\right] \geq \mathbb{E}\left[\frac{1}{T} \sum_t R^*(s_t) | \pi^*\right] - \epsilon.$$
  - Note:  $R^*$  is unknown.

# Solutions

---

- Ideally we want dense in time rewards to closely guide the agent closely along the way. Who will supply those shaped rewards?
  - We will manually design them: “cost function design by hand remains one of the ‘black arts’ of mobile robotics, and has been applied to untold numbers of robotic systems”
  - We will learn them from demonstrations: “rather than having a human expert tune a system to achieve desired behavior, the expert can demonstrate desired behavior and the robot can tune itself to match the demonstration”

# Solution – Learning from Demonstration

- Learning from demonstrations a.k.a. Imitation Learning: Supervision through an expert (teacher) that provides a set of demonstration trajectories: sequences of states and actions.
- **Imitation learning** is useful when it is easier for the expert to demonstrate the desired behavior rather than:
  - coming up with a reward function that would generate such behavior
  - coding up with the desired policy directly



# Solution – Learning from Demonstration

---

- Two broad approaches :
  - Direct: Supervised training of policy (mapping states to actions) using the demonstration trajectories as ground-truth (a.k.a. behavior cloning; especial case: it can be interactive, i.e., active behaviour cloning)
    - Can we directly learn the teacher's policy using supervised learning?
  - Indirect: Learn the unknown reward function/goal of the teacher, and derive the policy from these, a.k.a. Inverse Reinforcement Learning:
    - Inverse RL: Can we recover  $R$ ?
    - Apprenticeship learning via inverse RL: Can we then use this  $R$  to find a good policy?

# Solution – Learning from Demonstration

---

- Experts can be:
  - Humans
  - Optimal or near Optimal Planners/Controllers

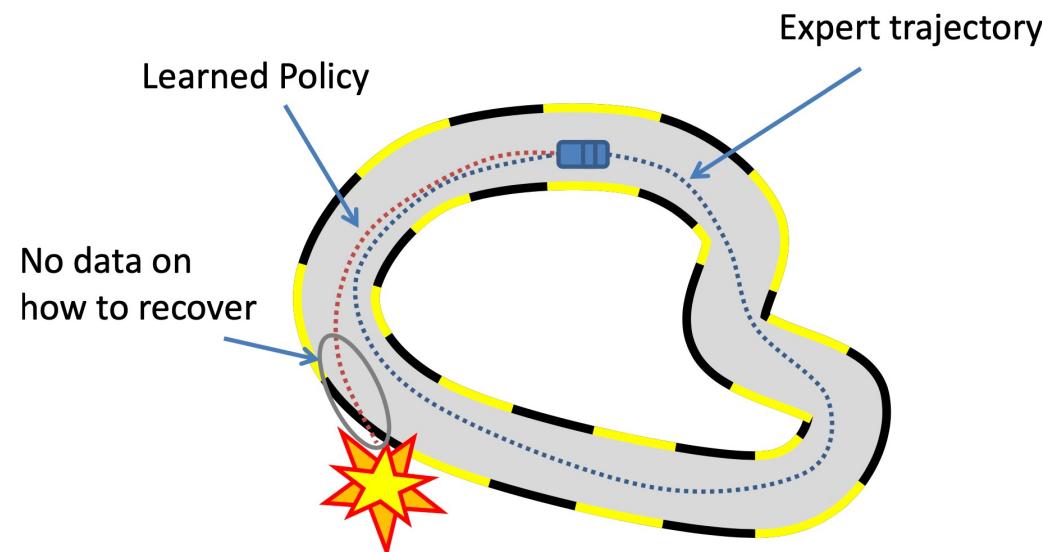
# Behavioral cloning

---

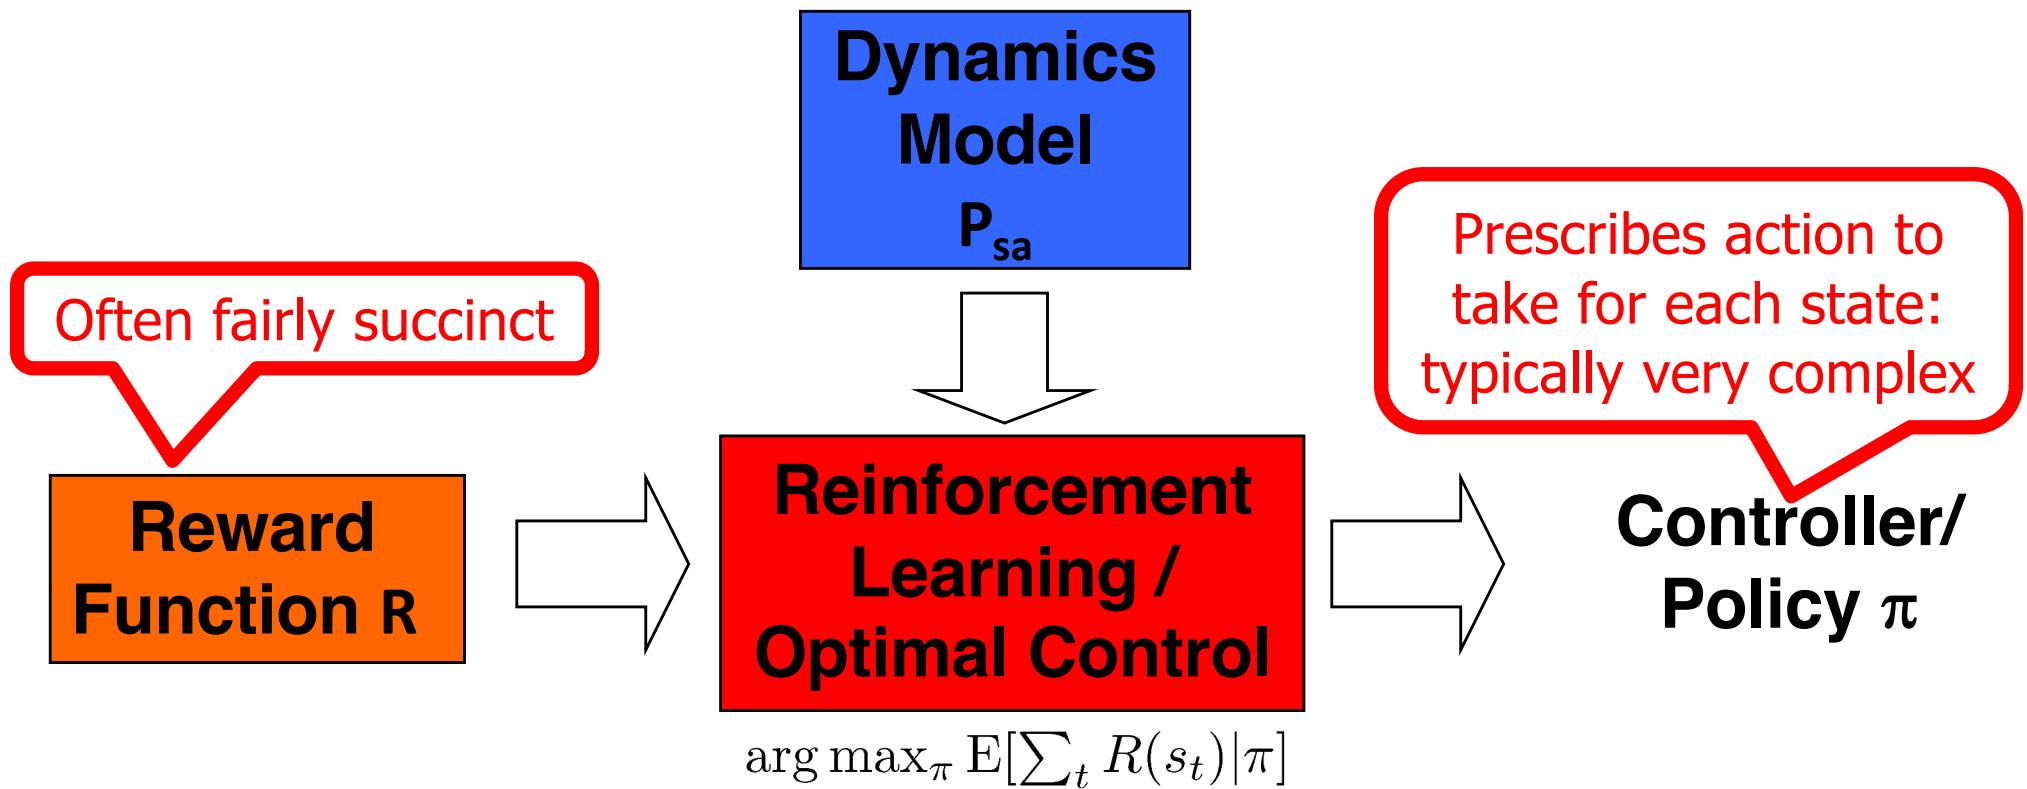
- Formulate as standard machine learning problem
  - Fix a policy class
    - E.g., support vector machine, neural network, decision tree, deep belief net, ...
    - Estimate a policy from the training examples  $(s_0, a_0), (s_1, a_1), (s_2, a_2), \dots$
- E.g.: <http://robotwhisperer.org/bird-muri/>
  - <https://youtu.be/hNsP6-K3Hn4>
- E.g., Pomerleau, 1989; Sammut et al., 1992; Kuniyoshi et al., 1994; Demiris & Hayes, 1994; Amit & Mataric, 2002.

# Behavioral cloning

- Limitations:
  - Underlying assumption: policy simplicity
  - makes mistakes: enters new states from which it cannot recover



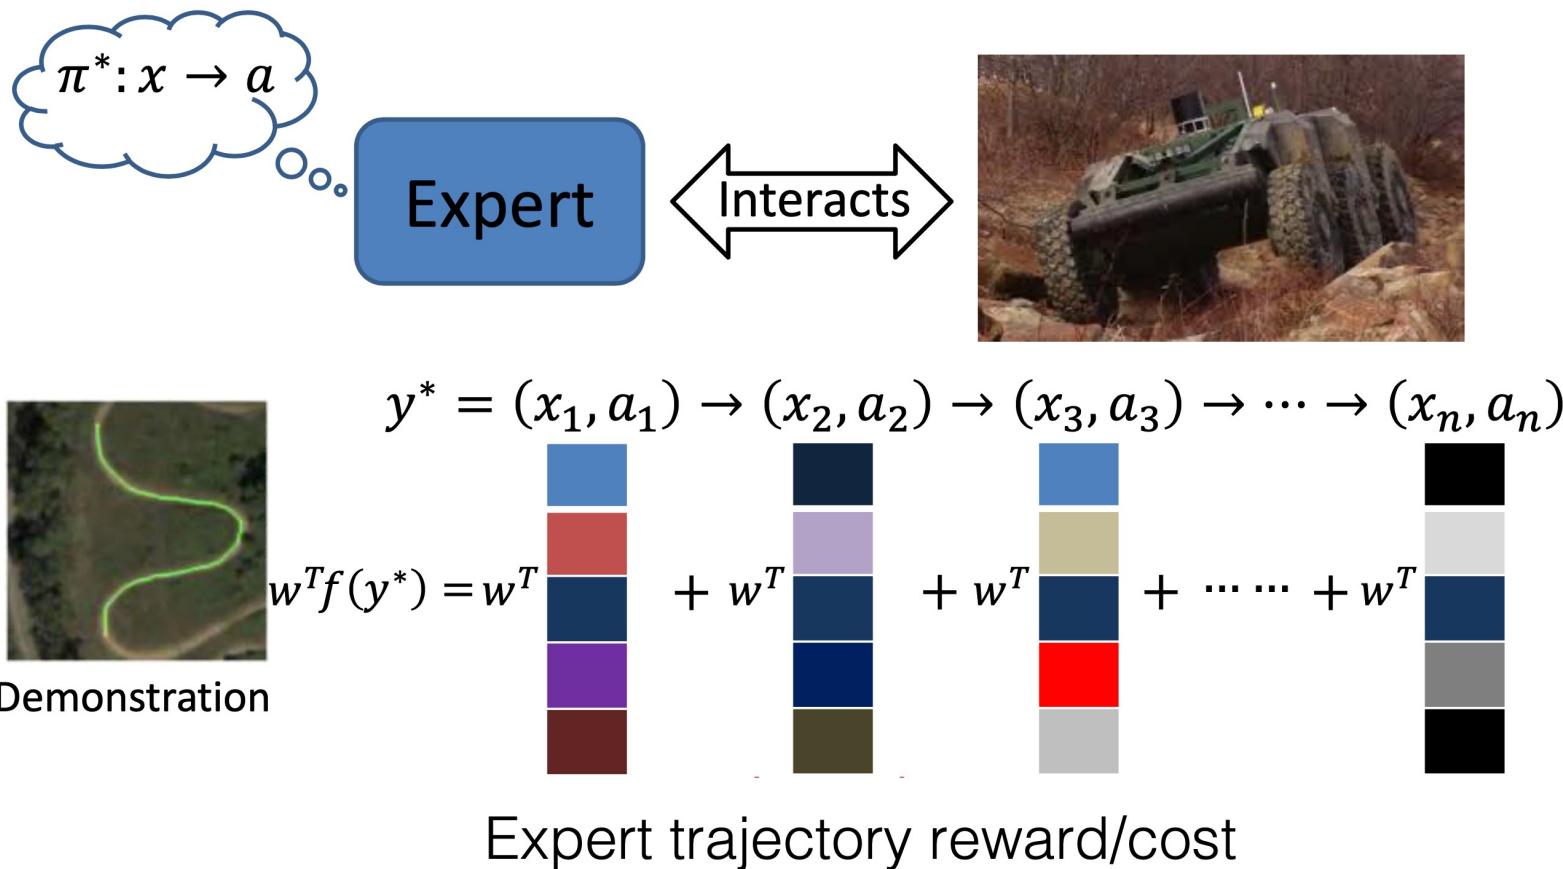
# Inverse RL



E.g.,  $R^* = w_1^* \mathbf{1}\{"\text{in right lane}"\} + w_2^* \mathbf{1}\{"\text{safe distance}"\}$

# Inverse RL

- Assumption: reward as a linear function of state features



# Inverse RL

- Assumption: reward as a linear function of state features

Find a reward function  $R^*$  which explains the expert behavior

i.e., assume expert follows optimal policy, given her  $R^*$

Find  $R^*$  such that

$$\mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t R^*(s_t) | \pi^*\right] \geq \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t R^*(s_t) | \pi\right] \quad \forall \pi$$

# Inverse RL

- Idea:

given:

states  $\mathbf{s} \in \mathcal{S}$ , actions  $\mathbf{a} \in \mathcal{A}$

(sometimes) transitions  $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$

samples  $\{\tau_i\}$  sampled from  $\pi^*(\tau)$

learn  $r_\psi(\mathbf{s}, \mathbf{a})$

reward parameters

...and then use it to learn  $\pi^*(\mathbf{a}|\mathbf{s})$

# RL vs. Inverse RL

given:

states  $\mathbf{s} \in \mathcal{S}$ , actions  $\mathbf{a} \in \mathcal{A}$

(sometimes) transitions  $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$

reward function  $r(\mathbf{s}, \mathbf{a})$

learn  $\pi^*(\mathbf{a}|\mathbf{s})$

given:

states  $\mathbf{s} \in \mathcal{S}$ , actions  $\mathbf{a} \in \mathcal{A}$

(sometimes) transitions  $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$

samples  $\{\tau_i\}$  sampled from  $\pi^*(\tau)$

learn  $r_\psi(\mathbf{s}, \mathbf{a})$

reward parameters

...and then use it to learn  $\pi^*(\mathbf{a}|\mathbf{s})$

# Inverse RL

---

- Idea:

1. Guess an initial reward function  $R(s)$
2. Learn policy  $\pi(s)$  that optimizes  $R(s)$
3. Whenever  $\pi(s)$  chooses action different from expert  $\pi^*(s)$ 
  - Update estimate of  $R(s)$  to assure value of  $\pi^*(s) >$  value of  $\pi(s)$
4. Go to 2

# Apprenticeship learning [Abbeel & Ng, 2004]

- Assume  $R_w(s) = w^\top \phi(s)$  for a feature map  $\phi : S \rightarrow \mathbb{R}^n$ .
- Initialize: pick some controller  $\pi_0$ .
- Iterate for  $i = 1, 2, \dots$  :
  - **“Guess” the reward function:**

Learning through reward functions rather than directly learning policies.

Find a reward function such that the teacher maximally outperforms all previously found controllers.

$$\max_{\gamma, w: \|w\|_2 \leq 1} \gamma$$

$$\text{s.t. } \mathbb{E} \left[ \sum_{t=0}^T R_w(s_t) | \pi^* \right] \geq \mathbb{E} \left[ \sum_{t=0}^T R_w(s_t) | \pi \right] + \gamma \quad \forall \pi \in \{\pi_0, \pi_1, \dots, \pi_{i-1}\}$$

- **Find optimal control policy**  $\pi_i$  for the current guess of the reward function  $R_w$ .
- If  $\gamma \leq \varepsilon/2$  exit the algorithm.

There is no reward function for which the teacher significantly outperforms thus-far found policies.

# Apprenticeship learning [Abbeel & Ng, 2004]

## ApprenticeshipLearning(*trajectories*)

$i$  indexes features,  $j$  indexes trajectories,  $k$  indexes policies

$k \leftarrow 0$

Initialize  $\pi_k$  at random

Repeat

$$\text{margin} = \max_w \min_{j,k} \sum_{i,t} w_i \phi_i(s_t^j, a_t^j) - \sum_{i,t} w_i \phi_i(s_t^j, \pi_k(s_t^j))$$

where  $w^*$  is the maximizing  $w$

$k \leftarrow k + 1$

$\pi_k \leftarrow$  optimal policy for  $w^*$  found by RL

Until  $\text{margin}$  is small enough

Return  $\pi_k$

# Apprenticeship learning [Abbeel & Ng, 2004] - Theoretical guarantees

## Theorem.

To ensure with probability at least  $1 - \delta$  that our algorithm returns a policy  $\pi$  such that

$$\mathbb{E}\left[\frac{1}{T} \sum_t R_w^*(s_t) | \pi\right] \geq \mathbb{E}\left[\frac{1}{T} \sum_t R_w^*(s_t) | \pi^*\right] - \epsilon.$$

it suffices that

we run for  $\frac{4n}{\epsilon^2}$  iterations,

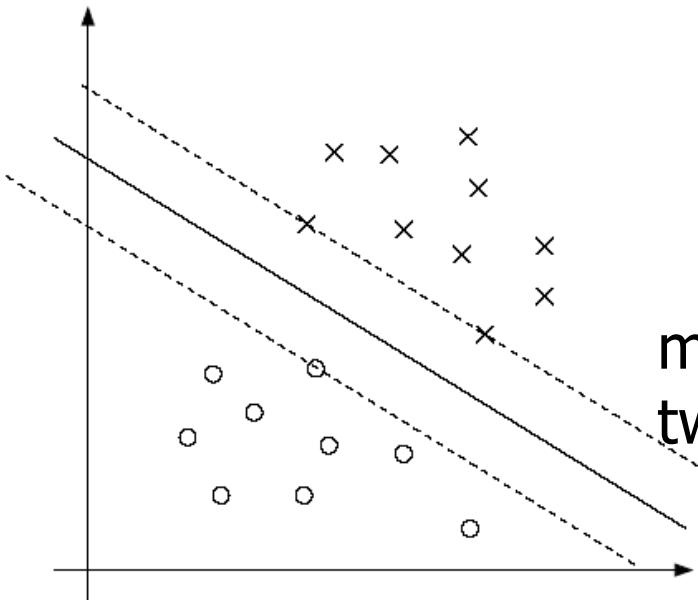
we have  $m \geq \frac{2n}{\epsilon^2} \log \frac{2n}{\delta}$  demonstrations.

- Guarantee w.r.t. unrecoverable reward function of teacher.
- Sample complexity does *not* depend on complexity of teacher's policy  $\pi^*$ .

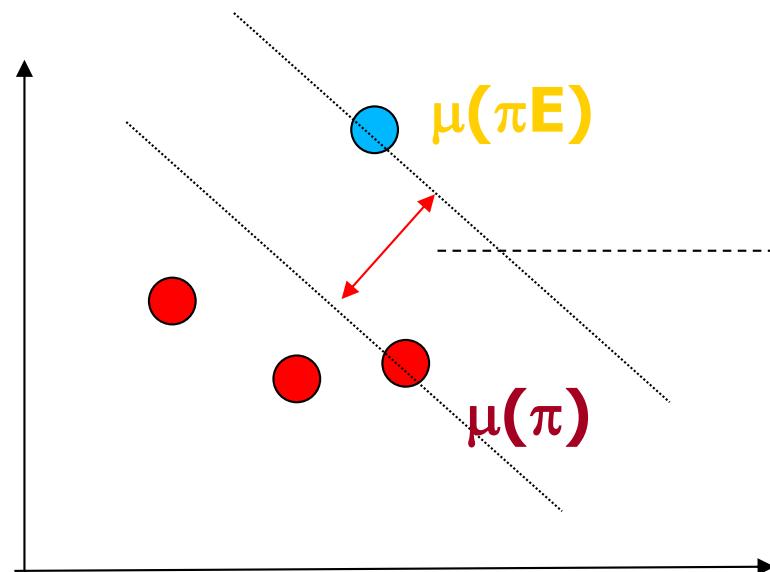
# Apprenticeship learning [Abbeel & Ng, 2004]

- For  $t = 1, 2, \dots$ 
  - **Inverse RL step:**
    - Estimate expert's reward function  $R(s) = w^T \phi(s)$  such that under  $R(s)$  the expert performs better than all previously found policies  $\{\pi_i\}$ .
  - **RL step:**
    - Compute optimal policy  $\pi_t$  for the estimated reward  $w$ .

# Apprenticeship learning [Abbeel & Ng, 2004]

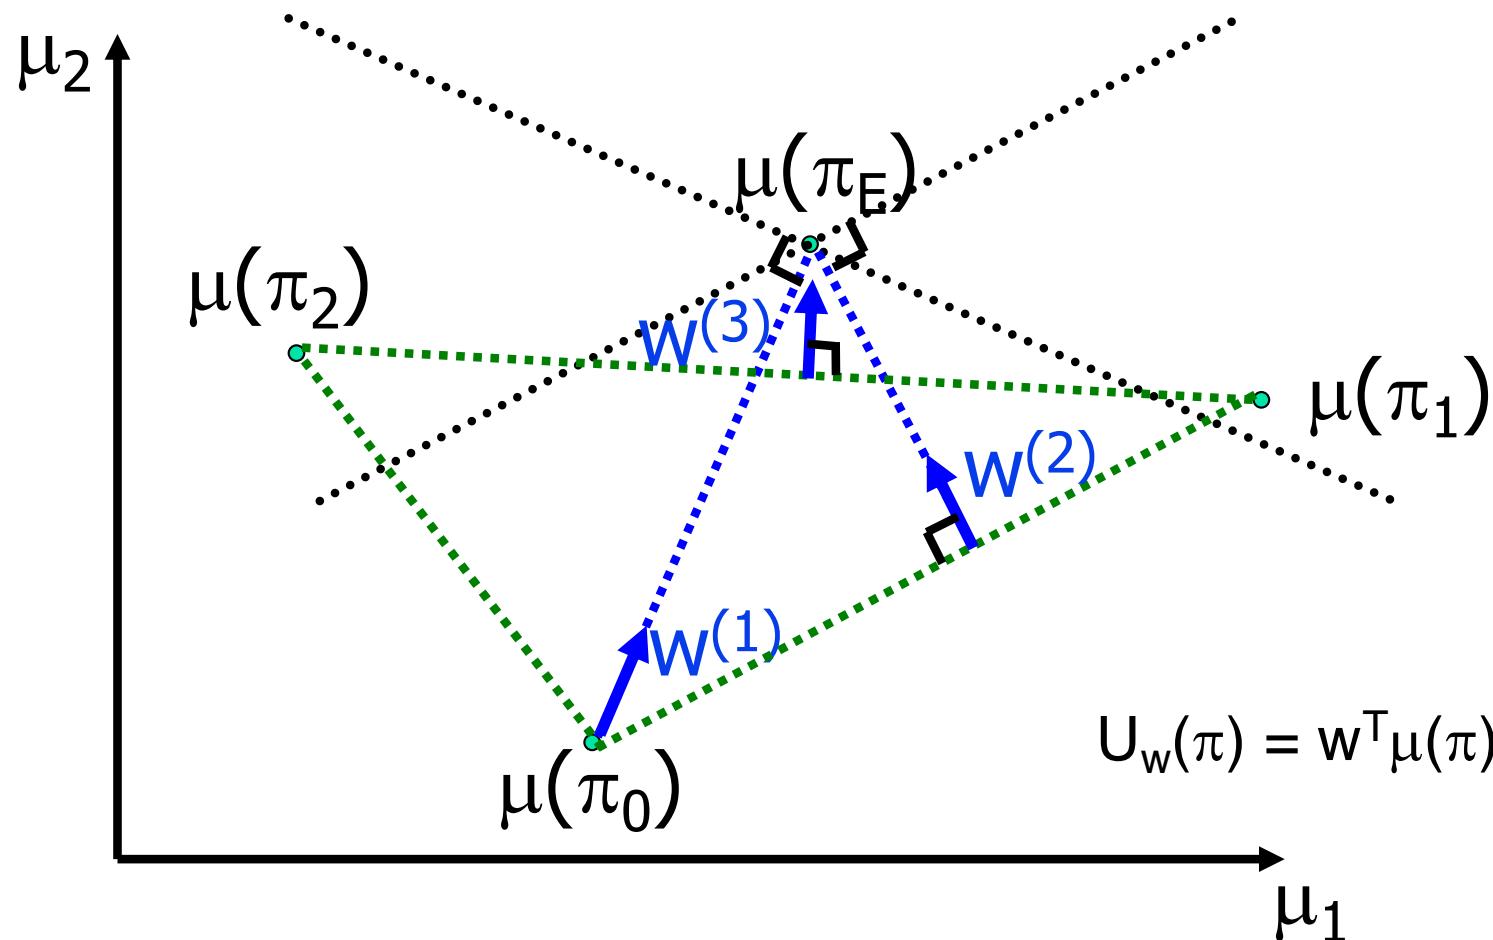


maximum margin hyperplane separating two sets of points



$|w^* \cdot \mu(\pi_E) - w^* \cdot \mu(\pi)|$   
 $= |Vw^*(\pi_E) - Vw^*(\pi)|$   
= maximal difference  
between expert policy's value  
function and 2<sup>nd</sup> to the  
optimal policy's value function

# Apprenticeship learning [Abbeel & Ng, 2004]



# Gridworld Experiment

- 128 x 128 grid world divided into 64 regions, each of size 16 x 16 ("macrocells").
- A small number of macrocells have positive rewards.
- For each macrocell, there is one feature  $\Phi_i(s)$  indicating whether that state  $s$  is in macrocell  $i$
- *Algorithm was also run on the subset of features  $\Phi_i(s)$  that correspond to non-zero rewards.*

# Gridworld Results

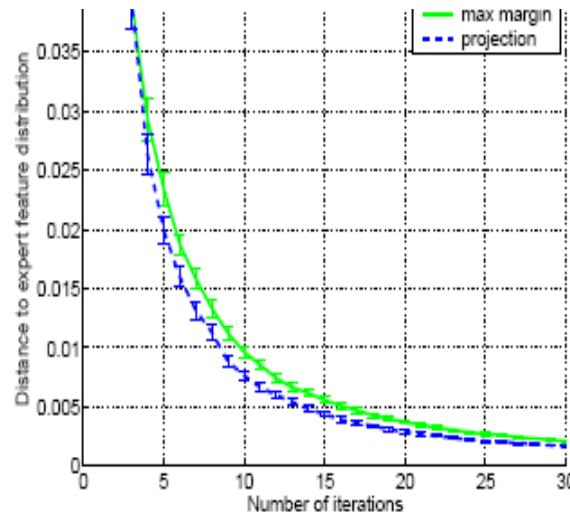
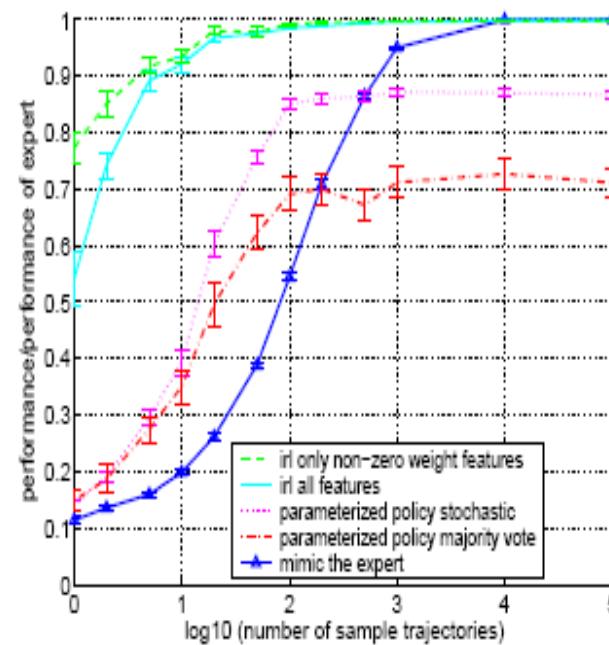


Figure 3. A comparison of the convergence speeds of the max-margin and projection versions of the algorithm on a 128x128 grid. Euclidean distance to the expert's feature expectations is plotted as a function of the number of iterations. We rescaled the feature expectations by  $(1 - \gamma)$  such that they are in  $[0, 1]^k$ . The plot shows averages over 40 runs, with 1 s.e. errorbars.

Distance to expert vs. #  
Iterations



Performance vs. #  
Trajectories

# Car Driving Experiment

- No explicit reward function at all!
- Expert demonstrates proper policy via 2 min. of driving time on simulator (1200 data points).
- 5 different “driver types” tried.
- Features: which lane the car is in, distance to closest car in current lane.
- Algorithm run for 30 iterations, policy hand-picked.
- Movie Time! (Expert left, IRL right)

# Examples

---

- **Learning for Control from Multiple Demonstrations** Adam Coates, Pieter Abbeel, Andrew Ng, ICML 2008
- **An Application of Reinforcement Learning to Aerobatic Helicopter Flight** Pieter Abbeel, Adam Coates, Morgan Quigley, Andrew Y. Ng, NIPS 2006
- <https://www.youtube.com/watch?v=0JL04JJjocc>

# Examples

- **Planning-based Prediction for Pedestrians** Brian Ziebart et al., IROS 2009
- <https://www.youtube.com/watch?v=hjOteEd7qwE>



# Examples

- **Data Driven Ghosting using Deep Imitation Learning** Hoang M. Le et al., SSAC 2017
- <https://www.youtube.com/watch?v=WI-WL2cj0CA>

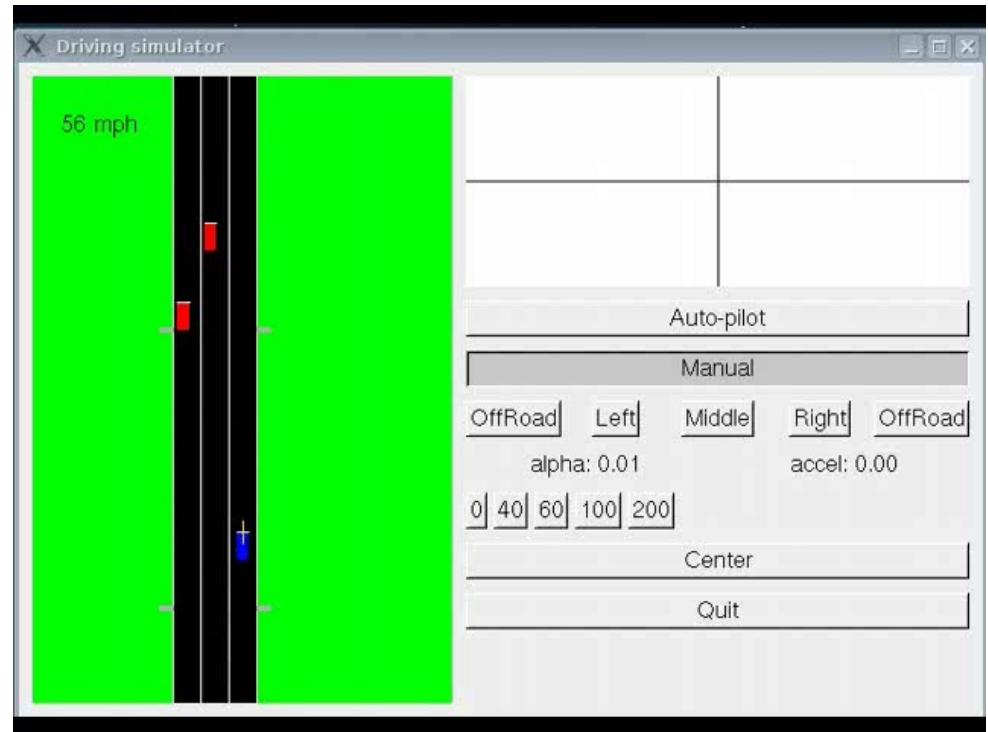


# Example: Highway driving

Teacher in Training World



Learned Policy in Testing World



- Input:
  - Dynamics model / Simulator  $P_{sa}(s_{t+1} | s_t, a_t)$
  - Teacher's demonstration: 1 minute in "training world"
  - Note:  $R^*$  is unknown.
  - Reward features: 5 features corresponding to lanes/shoulders; 10 features corresponding to presence of other car in current lane at different distances