

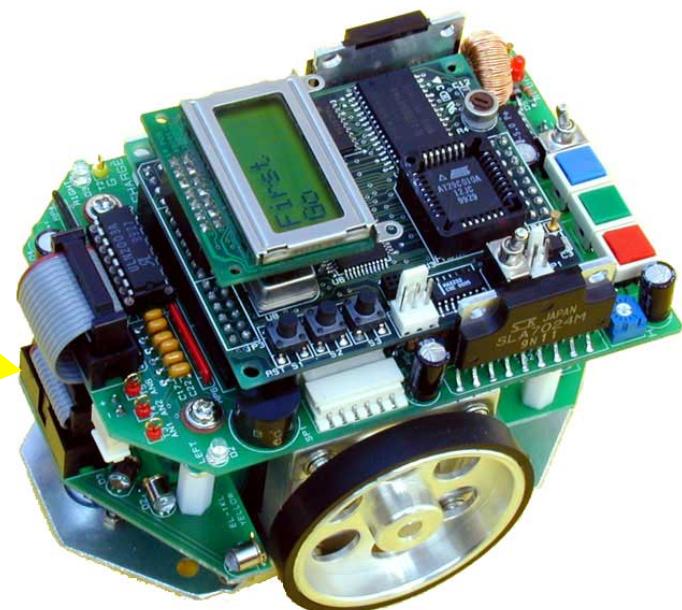
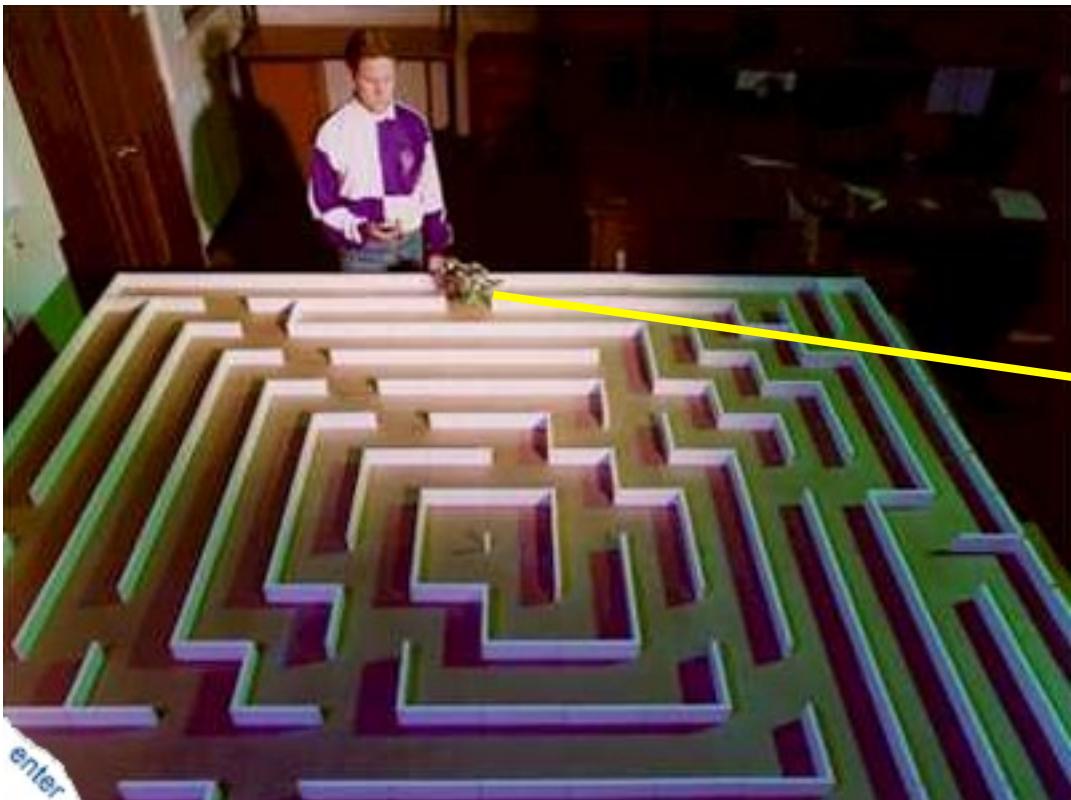
# Reinforcement Learning

Luís Macedo

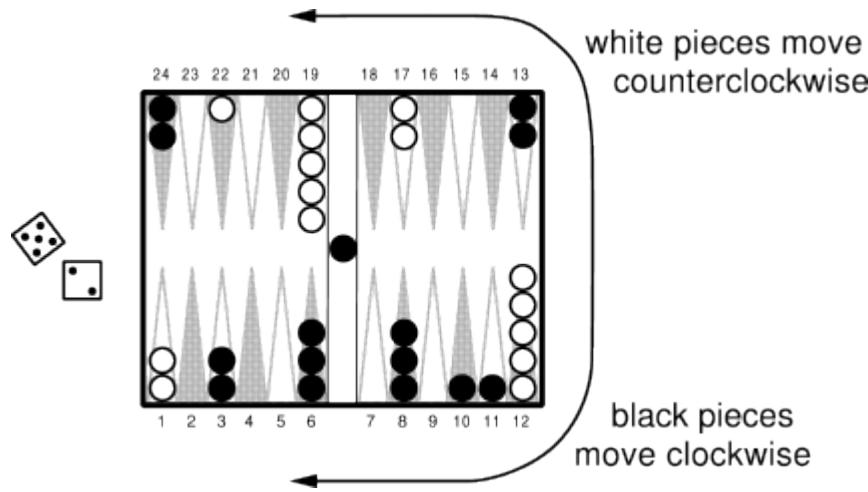
# So far ....

- Given an MDP model we know how to find optimal policies
  - ▲ Value Iteration or Policy Iteration
- But what if we don't have any form of model
  - ▲ In a Maze
  - ▲ All we can do is wander around the world observing what happens, getting rewarded and punished
- Enters reinforcement learning

# Micromouse



# Reinforcement Learning - Example



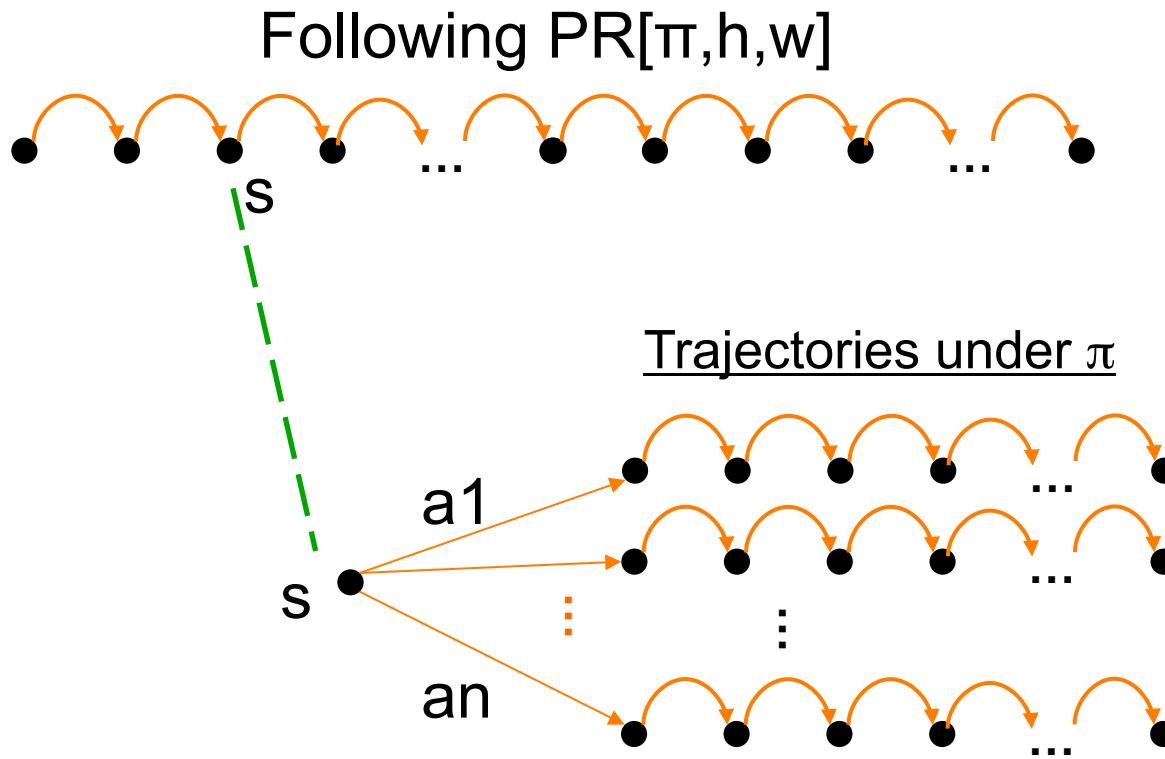
$(s_1, a_1), (s_2, a_2), \dots, (s_n, a_n)$  - Win

$$V(s) \leftarrow V(s) + \alpha(R(s) + \beta V(s') - V(s))$$

$(s_1, a_1), (s_2, a_2), \dots, (s_m, a_m)$  - Loss

What is the difference between TD and Rollout?

# Policy Rollout: Time Complexity



- To compute  $\text{PR}[\pi, h, w](s)$  for each action we need to compute  $w$  trajectories of length  $h$ 
  - Total of  $|A|hw$  calls to the simulator

# Markov Decision Processes - Recitation

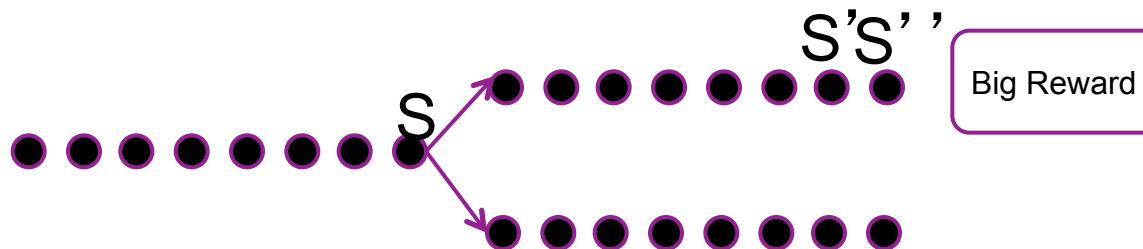
- $M = (S, A, T, R)$ 
  - ▲ S: States
  - ▲ A: Actions
  - ▲ T: Transition Probability
  - ▲ R: Reward
- Goal of the MDP?
  - ▲ Finding the Policy
    - Policy is the mapping from state to action
  - ▲ Two well known objectives
    - Average Reward
    - Discounted Reward

# Reinforcement Learning

- No knowledge of environment
  - ▲ Can only act in the world and observe states and reward
- Many factors make RL difficult:
  - ▲ Actions have **non-deterministic effects**
    - Which are initially unknown
  - ▲ **Rewards / punishments** are infrequent
    - Often at the end of long sequences of actions
    - How do we determine what action(s) were really responsible for reward or punishment?  
(credit assignment)
  - ▲ World is large and complex
- Nevertheless learner **must decide** what actions to take
  - ▲ We will assume the world behaves as an MDP

# Delayed Reward Problem

- Delayed Reward Make it hard to learn
- The choice in State S was important but, it seems the action in S' lead to the big reward in S''
- How you deal with this problem?



- How about the performance comparison of VI and PI in this example?

# Reinforcement Learning

- Something is unknown
- Learning and Planning at the same time
- Ultimate learning and planning paradigm
- Scalability is a big issue, Very Challenging!
- Ironically, RL was most successful in Real Application even more than STRIPS Planning!

- ▲ Zhang, W., Dietterich, T. G., (1995). **A Reinforcement Learning Approach to Job-shop Scheduling**
- ▲ G. Tesauro (1994). "TD-Gammon, A Self-Teaching Backgammon Program Achieves Master-level Play" in Neural Computation
- ▲ Reinforcement Learning for Vulnerability Assessment in Peer-to-Peer Networks, IAAI 2008
  - Policy Gradient Update

# Two Key Aspect in RL

- How we update the value function or policy?
  - ▲ How do we form training data
  - ▲ Sequence of  $(s, a, r)$ ....
- How we explore?
  - ▲ Exploit or Exploration

# Category of Reinforcement Learning

- Model-based RL
  - ▲ Constructs domain transition model, MDP
    - E<sup>3</sup> – Kearns and Singh
- Model-free RL
  - ▲ Only concerns policy
    - Q-Learning - Watkins
- Active Learning (Off-Policy Learning)
  - ▲ Q-Learning
- Passive Learning (On-Policy learning)
  - ▲ Sarsa - Sutton

# Passive RL

**”If you study hard, you will be blessed”**  
**“what is the value of studying hard?”**  
**“You will learn from RL”**

# Policy Evaluation

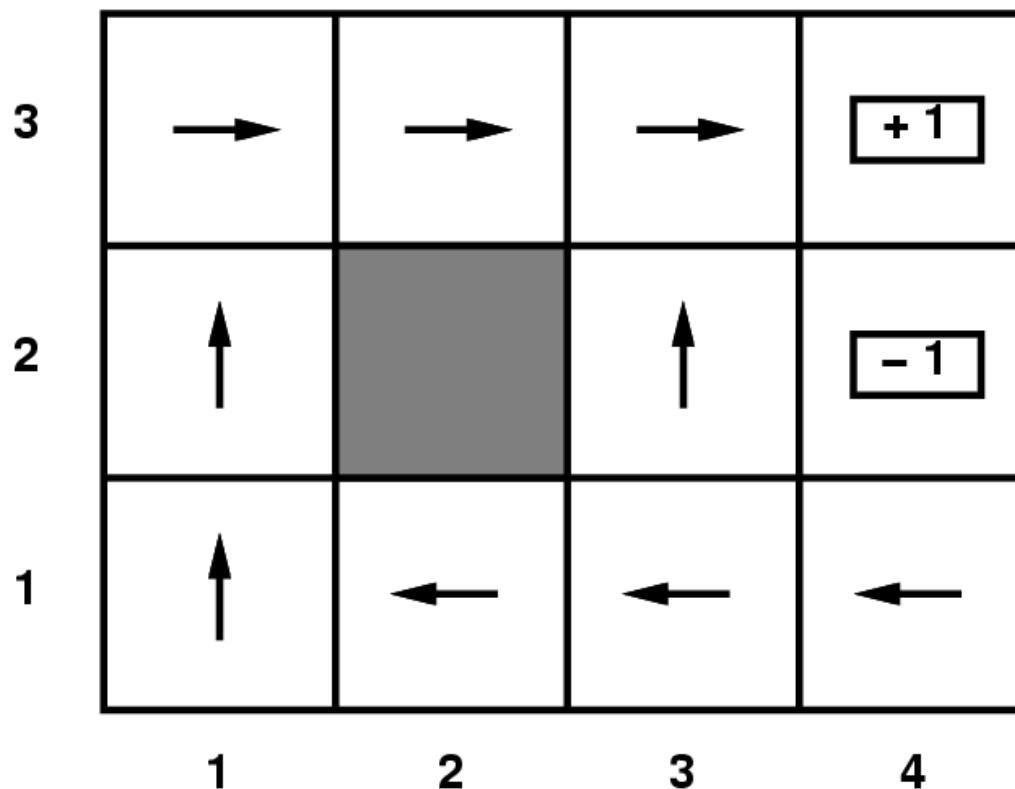
- Remember the formula?

$$V_{\pi}(s) = R(s) + \beta \sum_{s'} T(s, \pi(s), s') \cdot V_{\pi}(s')$$

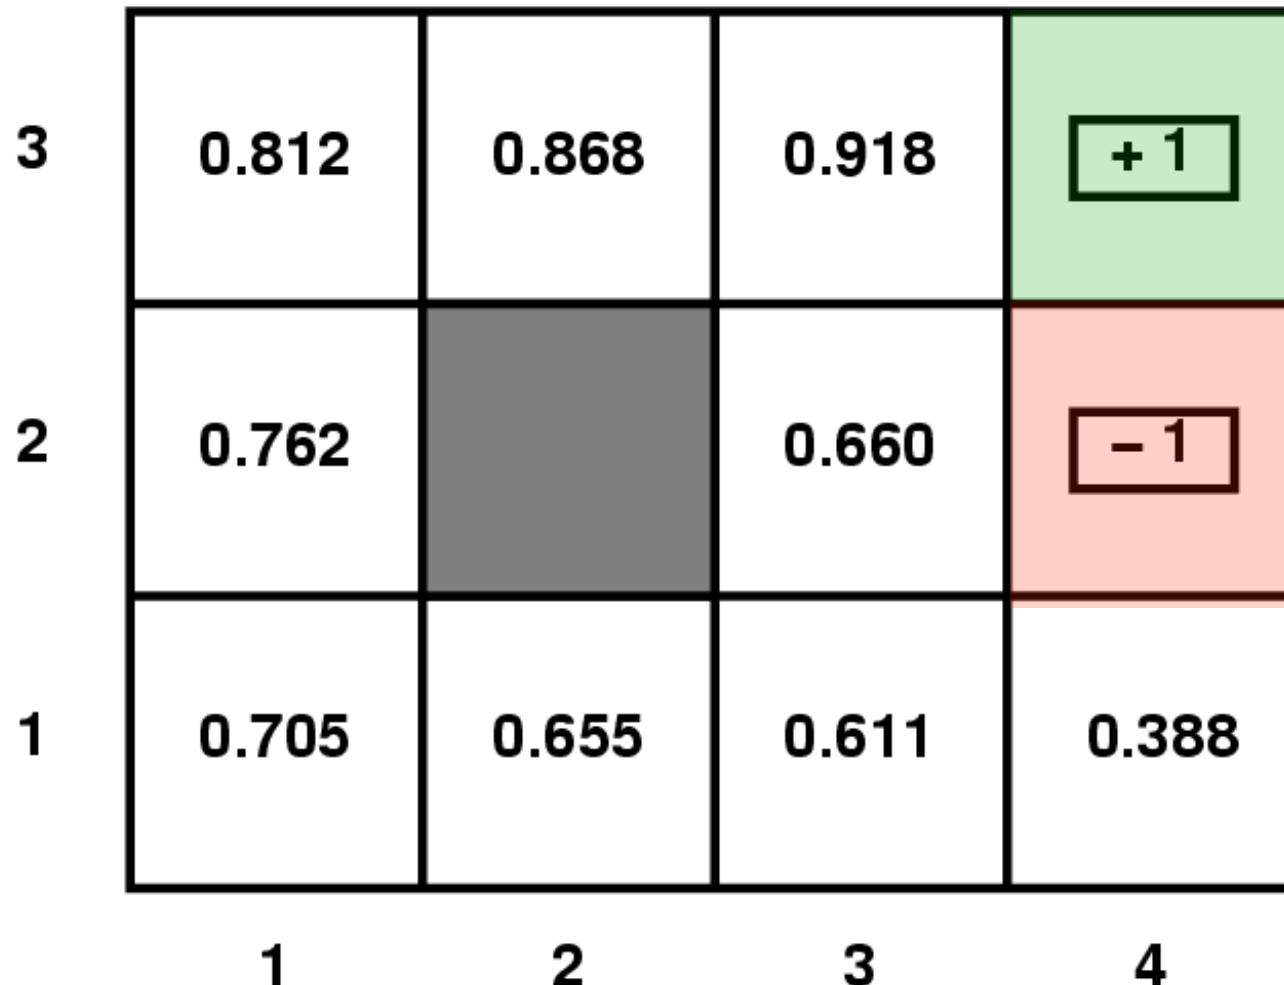
- Can we do that with RL?
  - ▲ What is missing?
  - ▲ What needs to be done?
- What do we do after policy evaluation?
  - ▲ Policy Update

# Example: Passive RL

- Suppose given policy
- Want to determine how good it is

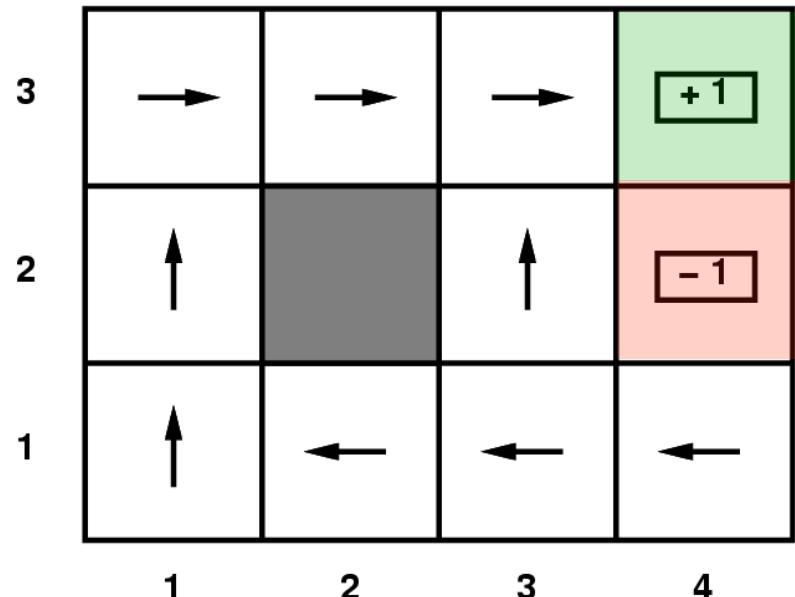


# Objective: Value Function



# Passive RL

- Given policy  $\pi$ ,
  - ▲ estimate  $V^\pi(s)$
- Not given
  - ▲ transition matrix, nor
  - ▲ reward function!
- Simply follow the policy for many epochs
- Epochs: training sequences



$(1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (3,3) \rightarrow (3,4) \underline{+1}$   
 $(1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (3,3) \rightarrow (3,2) \rightarrow (3,3) \rightarrow (3,4) \underline{+1}$   
 $(1,1) \rightarrow (2,1) \rightarrow (3,1) \rightarrow (3,2) \rightarrow (4,2) \underline{-1}$

# Approach 1: Direct Estimation

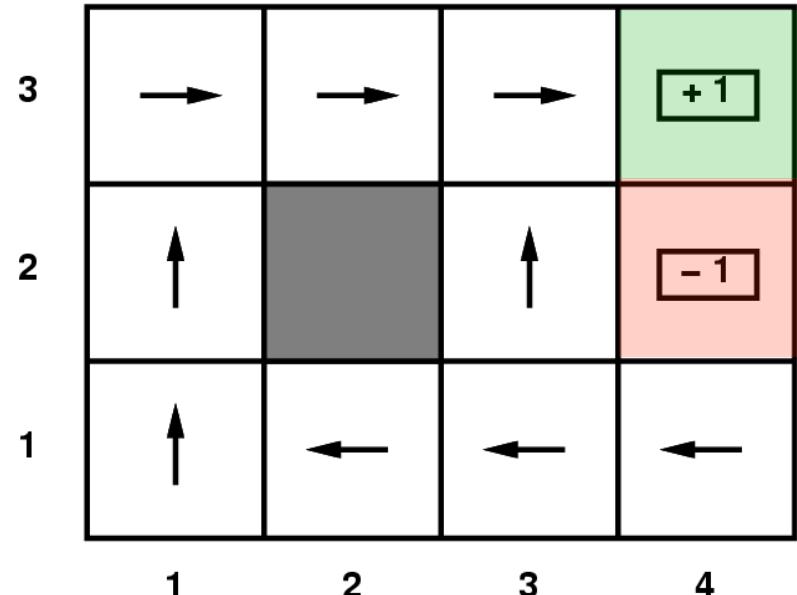
- Direct estimation (model free)
  - ▲ Estimate  $V^\pi(s)$  as average total reward of epochs containing s (calculating from s to end of epoch)
- **Reward to go** of a state s

the sum of the (discounted) rewards from that state until a terminal state is reached

- Key: use observed **reward to go** of the state as the direct evidence of the actual expected utility of that state
- Averaging the reward to go samples will converge to true value at state

# Passive RL

- Given policy  $\pi$ ,
  - ▲ estimate  $V^\pi(s)$
- Not given
  - ▲ transition matrix, nor
  - ▲ reward function!
- Simply follow the policy for many epochs
- Epochs: training sequences



$(1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (3,3) \rightarrow (3,4) \underline{+1}$   
 $\underline{\quad 0.57 \quad 0.64 \quad 0.72 \quad 0.81 \quad 0.9 \quad}$

$(1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (3,3) \rightarrow (3,2) \rightarrow (3,3) \rightarrow (3,4) \underline{+1}$

$(1,1) \rightarrow (2,1) \rightarrow (3,1) \rightarrow (3,2) \rightarrow (4,2) \underline{-1}$

# Direct Estimation

- Converge very slowly to correct utilities values (requires a lot of sequences)
- Doesn't exploit Bellman constraints on policy values

$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

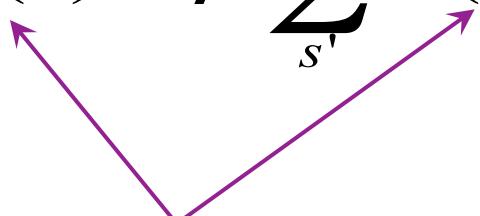
How can we incorporate constraints?

## Approach 2: Adaptive Dynamic Programming (ADP)

- ADP is a model based approach
  - ▲ Follow the policy for awhile
  - ▲ Estimate transition model based on observations
  - ▲ Learn reward function
  - ▲ Use estimated model to compute utility of policy

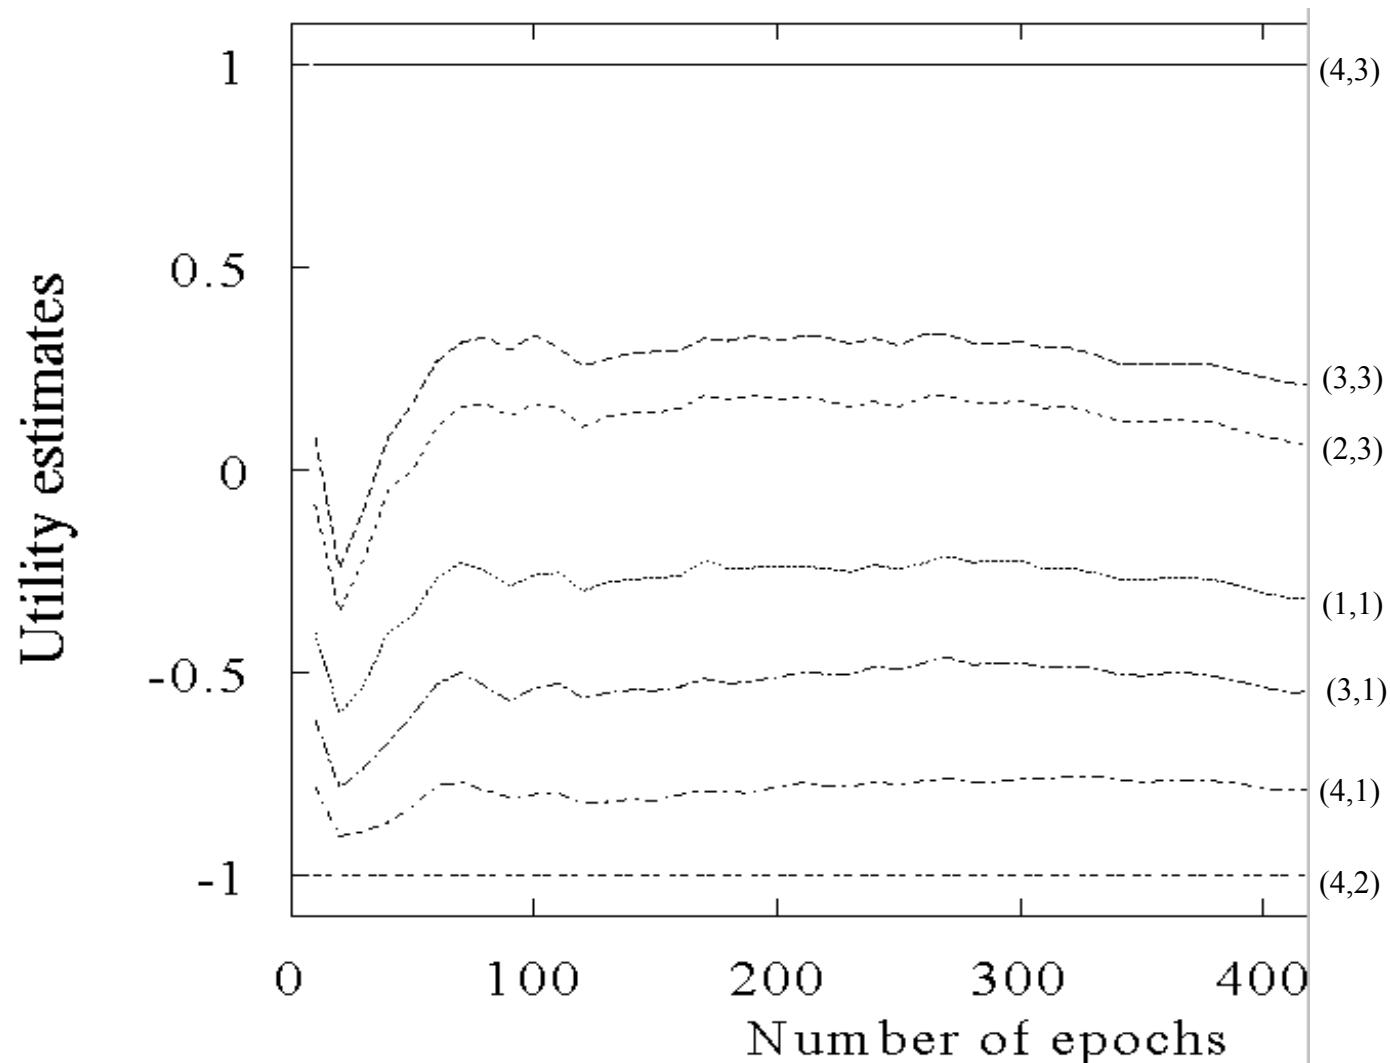
$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

learned



- How can we estimate transition model  $T(s, a, s')$ ?
  - ▲ Simply the fraction of times we see  $s'$  after taking  $a$  in state  $s$ .
  - ▲ NOTE: Can bound error with Chernoff bounds if we want (will see Chernoff bound later in course)

# ADP learning curves



# Approach 3: Temporal Difference Learning (TD)

- Can we avoid the computational expense of full DP policy evaluation?
  - Temporal Difference Learning
    - ▲ Do local updates of utility/value function on a **per-action** basis
    - ▲ Don't try to estimate entire transition function!
    - ▲ For each transition from  $s$  to  $s'$ , we perform the following update:

- Intuitively moves us closer to satisfying Bellman constraint

$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

## Why?

# Aside: Online Mean Estimation

- Suppose that we want to incrementally compute the mean of a sequence of numbers
    - ▲ E.g. to estimate the expected value of a random variable from a sequence of samples.

$$\hat{X}_{n+1} = \frac{1}{n+1} \sum_{i=1}^{n+1} x_i = \frac{1}{n} \sum_{i=1}^n x_i + \frac{1}{n+1} \left( x_{n+1} - \frac{1}{n} \sum_{i=1}^n x_i \right)$$

$$= \hat{X}_n + \frac{1}{n+1} (x_{n+1} - \hat{X}_n)$$

↑  
 average of n+1 samples  
↑  
 learning rate  
↑  
 sample n+1

- Given a new sample  $x(n+1)$ , the new mean is the old estimate (for  $n$  samples) plus the weighted difference between the new sample and old estimate

## Approach 3: Temporal Difference Learning (TD)

- TD update for transition from  $s$  to  $s'$ :

$$V^\pi(s) = V^\pi(s) + \alpha(R(s) + \beta V^\pi(s') - V^\pi(s))$$

learning rate      (noisy) sample of utility  
based on next state

- So the update is maintaining a “mean” of the (noisy) utility samples
- If the learning rate decreases with the number of samples (e.g.  $1/n$ ) then the utility estimates will converge to true values!

$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

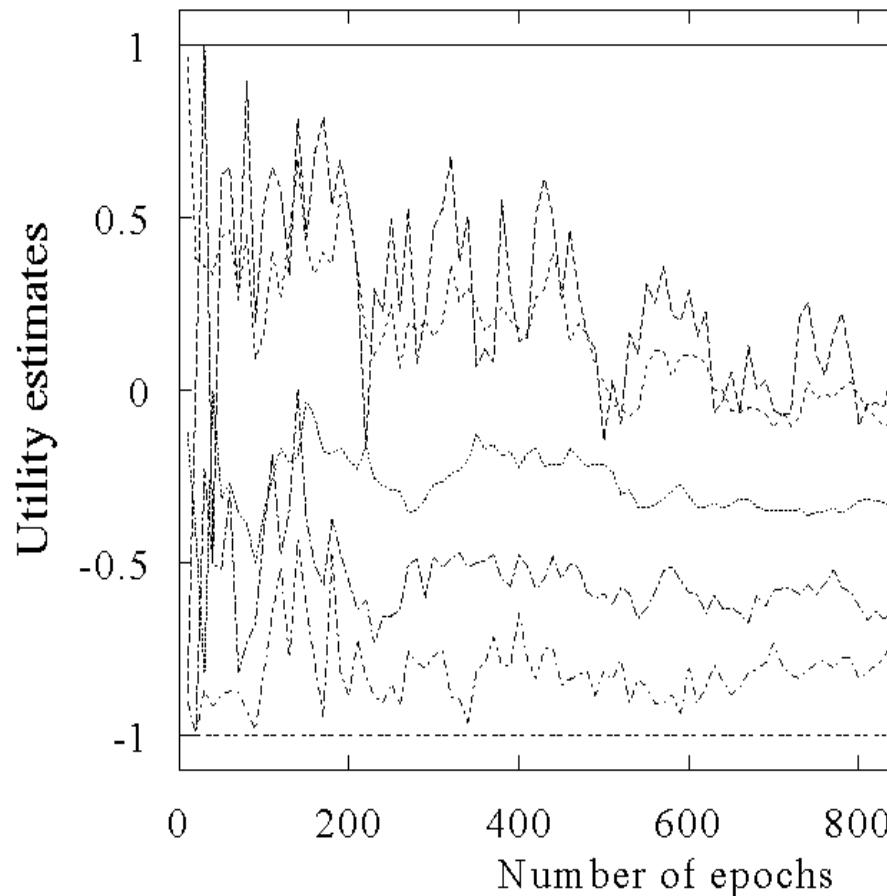
# Approach 3: Temporal Difference Learning (TD)

- TD update for transition from  $s$  to  $s'$ :

- When  $V$  satisfies Bellman constraints then **expected update** is 0.

$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

# The TD learning curve



- **Tradeoff:** requires more training experience (epochs) than ADP but much less computation per epoch
- Choice depends on relative cost of experience vs. computation

# Comparisons

- Direct Estimation (model free)
  - ▲ Simple to implement
  - ▲ Each update is fast
  - ▲ Does not exploit Bellman constraints
  - ▲ Converges slowly
- Adaptive Dynamic Programming (model based)
  - ▲ Harder to implement
  - ▲ Each update is a full policy evaluation (expensive)
  - ▲ Fully exploits Bellman constraints
  - ▲ Fast convergence (in terms of updates)
- Temporal Difference Learning (model free)
  - ▲ Update speed and implementation similar to direct estimation
  - ▲ Partially exploits Bellman constraints---adjusts state to ‘agree’ with observed successor
    - Not **all** possible successors
  - ▲ Convergence in between direct estimation and ADP

# Active RL

**”If you study hard, you will be blessed”  
“what is the value of studying hard?”  
“You will learn from RL”**

**“Well, I will not study hard always, I will try some music  
playing career and see if I can be an American Idol” - Sanjaya**

# Active Reinforcement Learning

- So far, we've assumed agent *has* policy
  - ▲ We just try to learn how good it is
- Now, suppose agent must learn a good policy (ideally optimal)
  - ▲ While acting in uncertain world

# Naïve Approach

1. Act Randomly for a (long) time
  - ▲ Or systematically explore all possible actions
2. Learn
  - ▲ Transition function
  - ▲ Reward function
3. Use value iteration, policy iteration, ...
4. Follow resulting policy thereafter.

**Will this work?** Yes (if we do step 1 long enough)

**Any problems?** We will act randomly for a long time before exploiting what we know.

# Revision of Naïve Approach

1. Start with initial utility/value function and initial model
2. Take greedy action according to value function|  
(this requires using our estimated model to do “lookahead”)
3. Update estimated model
4. Perform step of ADP ;; update value function
5. Goto 2

This is just ADP but we follow the greedy policy suggested by current value estimate

Will this work? No. Gets stuck in local minima.

What can be done?

# Exploration versus Exploitation

- Two reasons to take an action in RL
  - ▲ **Exploitation**: To try to get reward. We exploit our current knowledge to get a payoff.
  - ▲ **Exploration**: Get more information about the world. How do we know if there is not a pot of gold around the corner.
- To explore we typically need to take actions that do not seem best according to our current model.
- Managing the trade-off between exploration and exploitation is a critical issue in RL
- Basic intuition behind most approaches:
  - ▲ Explore more when knowledge is weak
  - ▲ Exploit more as we gain knowledge

# ADP-based RL

1. Start with initial utility/value function
2. Take action according to an **explore/exploit policy**  
(explores more early on and gradually becomes greedy)
3. Update estimated model
4. Perform step of ADP
5. Goto 2

This is just ADP but we follow the explore/exploit policy

Will this work? Depends on the explore/exploit policy.

Any ideas?

# Explore/Exploit Policies

- Greedy action is action maximizing estimated Q-value

$$Q(s, a) = R(s) + \beta \sum_{s'} T(s, a, s') V(s')$$

- ▲ where V is current value function estimate, and R, T are current estimates of model
- ▲  $Q(s, a)$  is the expected value of taking action  $a$  in state  $s$  and then getting the estimated value  $V(s')$  of the next state  $s'$

- Want an exploration policy that is **greedy in the limit of infinite exploration (GLIE)**

- ▲ Guarantees convergence

- Solution 1

- ▲ On time step  $t$  select random action with probability  $p(t)$  and greedy action with probability  $1-p(t)$
- ▲  $p(t) = 1/t$  will lead to convergence, but is slow

# Explore/Exploit Policies

- Greedy action is action maximizing estimated Q-value

$$Q(s, a) = R(s) + \beta \sum_{s'} T(s, a, s') V(s')$$

► where V is current value function estimate, and R, T are current estimates of model

- Solution 2: Boltzmann Exploration
  - Select action a with probability,

$$\Pr(a | s) = \frac{\exp(Q(s, a) / T)}{\sum_{a' \in A} \exp(Q(s, a') / T)}$$

- T is the temperature. Large T means that each action has about the same probability. Small T leads to more greedy behavior.
- Typically start with large T and decrease with time

# Alternative Approach: Exploration Functions

1. Start with initial utility/value function
2. Take **greedy** action
3. Update estimated model
4. Perform step of optimistic ADP  
(uses exploration function in value iteration)  
(inflates value of actions leading to unexplored regions)
5. Goto 2

What do we mean by exploration function in VI?

# Exploration Functions

- Recall that value iteration iteratively performs the following update at all states:

$$V(s) \leftarrow R(s) + \beta \max_a \sum_{s'} T(s, a, s') V(s')$$

- We want the update to make actions that lead to unexplored regions look good
- Implemented by ***exploration function f(u,n)***:
  - assigning a higher utility estimate to relatively unexplored action state pairs
  - change the updating rule of value function to

$$V^+(s) \leftarrow R(s) + \beta \max_a f\left( \sum_{s'} T(s, a, s') V^+(s'), N(a, s) \right)$$

- $V^+$  denote the optimistic estimate of the utility
- $N(s,a)$  = number of times that action  $a$  was taken from state  $s$

What properties should  $f(u,n)$  have?

# Exploration Functions

$$V^+(s) \leftarrow R(s) + \beta \max_a f\left( \sum_{s'} T(s, a, s') V^+(s'), N(a, s) \right)$$

- Properties of  $f(u, n)$ ?
  - ▲ If  $n > N_e$        $u$     i.e. normal utility
  - ▲ Else,                   $R^+$  i.e. max possible value of a state
- The agent will behave initially as if there were wonderful rewards scattered all over around— optimistic .
- But after actions are tried enough times we will perform standard “non-optimistic” value iteration
- Note that all of these exploration approaches assume that exploration will not lead to unrecoverable disasters (falling off a cliff).

# TD-based Active RL

1. Start with initial utility/value function
2. Take action according to an **explore/exploit policy**  
(should converge to greedy policy, i.e. GLIE)
3. Update estimated model
4. Perform TD update

$$V(s) \leftarrow V(s) + \alpha(R(s) + \beta V(s') - V(s))$$

$V(s)$  is new estimate of optimal value function at state  $s$ .

5. Goto 2

Just like TD for passive RL, but we follow explore/exploit policy

Given the usual assumptions about learning rate and GLIE,  
TD will converge to an optimal value function!

# TD-based Active RL

1. Start with initial utility/value function
2. Take action according to an **explore/exploit policy**  
(should converge to greedy policy, i.e. GLIE)
3. Update estimated model
4. Perform TD update

$$V(s) \leftarrow V(s) + \alpha(R(s) + \beta V(s') - V(s))$$

$V(s)$  is new estimate of optimal value function at state  $s$ .

5. Goto 2

Requires an estimated model. Why?

To compute  $Q(s,a)$  for greedy policy execution

Can we construct a model-free variant?

# Q-Learning: Model-Free RL

- Instead of learning the optimal value function  $V$ , directly learn the optimal Q function.
  - Recall  $Q(s,a)$  is expected value of taking action  $a$  in state  $s$  and then following the optimal policy thereafter
- The optimal Q-function satisfies  $V(s) = \max_{a'} Q(s, a')$  which gives:

$$\begin{aligned} Q(s, a) &= R(s) + \beta \sum_{s'} T(s, a, s') V(s') \\ &= R(s) + \beta \sum_{s'} T(s, a, s') \max_{a'} Q(s, a') \end{aligned}$$

- Given the Q function we can act optimally by select action greedily according to  $Q(s,a)$

How can we learn the Q-function directly?

# Q-Learning: Model-Free RL

Bellman constraints on optimal Q-function:

$$Q(s, a) = R(s) + \beta \sum_{s'} T(s, a, s') \max_{a'} Q(s', a')$$

- We can perform updates after each action just like in TD.
  - ▲ After taking **action a** in **state s** and reaching **state s'** do:  
(note that we directly observe reward  $R(s)$ )

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \beta \max_{a'} Q(s', a') - Q(s, a))$$

  
(noisy) sample of Q-value  
based on next state

# Q-Learning

1. Start with initial Q-function (e.g. all zeros)
2. Take action according to an **explore/exploit policy** (should converge to greedy policy, i.e. GLIE)
3. Perform TD update

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \beta \max_{a'} Q(s', a') - Q(s, a))$$

$Q(s, a)$  is current estimate of optimal Q-function.

4. Goto 2

- Does not require model since we learn Q directly!
- Uses explicit  $|S| \times |A|$  table to represent Q
- Explore/exploit policy directly uses Q-values
  - E.g. use Boltzmann exploration.

# Direct Policy RL

- Why?
- Again, ironically, policy gradient based approaches were successful in many real applications
- Actually we do this.
  - ▲ From 10 Commandments
  - ▲ “You Shall Not Murder”
  - ▲ “Do not have any other gods before me”
- How we design an policy with parameters?
  - ▲ Multinomial distribution of actions in a state
  - ▲ Binary classification for each action in a state

# Policy Gradient Algorithm

- $J(\mathbf{w}) = E_{\mathbf{w}}[\sum_{t=0..T} \gamma^t c_t]$  (failure prob., makespan, ...)
  - minimise  $J$  by
    - ▲ computing gradient  $\nabla J(\mathbf{w}) = \left[ \frac{\partial J}{\partial \mathbf{w}_1}, \frac{\partial J}{\partial \mathbf{w}_2}, \dots, \frac{\partial J}{\partial \mathbf{w}_k} \right]$
    - ▲ stepping the parameters away  $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \nabla J(\mathbf{w})$
  - until convergence
- 
- Gradient Estimate [Sutton et.al.'99, Baxter & Bartlett'01]
  - Monte Carlo estimate from trace  $s_1, a_1, c_1, \dots, s_T, a_T, c_T$ 
    - ▲  $e_{t+1} = e_t + r_w \log \Pr(a_{t+1}|s_t, \mathbf{w}_t)$
    - ▲  $\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \gamma^t c_t e_{t+1}$
  - Successfully used in many applications!

# RL Summary

- Wondering around the world provides
  - ▲ Training Data
  - ▲ Each episode is a sample of the current policy
  - ▲ Sampling in RL is unique, the policy is stochastic and still provides a guarantee to convergence
  - ▲ Where to sample is very important!
    - Exploration vs. Exploitation